FirecREST Agentic Workbench

Natural-language HPC workflows, with guardrails and reproducibility

The problem

FirecREST gives HPC a REST API, but a researcher who is not an HPC expert still has to remember endpoints, job-script syntax and scheduler options, and has no cheap way to prove later that a result can be reproduced.

What we built

An agentic layer on top of FirecREST. The researcher asks in plain language (“run the matrix benchmark on 4 nodes and tell me when it’s done”); the agent does the rest.

Architecture

Request in natural language → agent (Hermes) → MCP gateway → **FirecREST MCP wrapper** → FirecREST → Slurm. Every FirecREST call goes through a single client module; no other component talks to the API.

What is new

- **Guardrailed templates:** jobs are launched from vetted templates with enforced limits.
- **Human approval** before anything is submitted.
- **Guardrail refusals:** dangerous or out-of-policy requests are blocked, and nothing is submitted.
- **Reproducibility card:** inputs, template, parameters and output hashes are recorded; a job can be relaunched from the card and its outputs verified by hash.

Try it

Everything runs against the official FirecREST demo stack; no real credentials are needed.

Repository: https://github.com/VadaLinux/firecrest-agentic-workbench

Repository QR code

FirecREST Hackathon, CSCS / ETH Zurich, 3 November 2026.